

Evidence lower bounds with built-in baselines

Justin

December 29, 2022

1 Problem setup

The goal is to maximize the log marginal likelihood,

$$\log p(x) = \log \sum_z p(x, z), \quad (1)$$

for a latent variable model. The derivative of the above is

$$\nabla \log \sum_z p(x, z) = \frac{p(x, z)}{p(x)} \nabla \log p(x, z) = p(z | x) \nabla \log p(x, z), \quad (2)$$

the expected gradient under the posterior. When exact marginalization is intractable, a common approach is to introduce a variational approximation to the posterior $q(z | x)$ and optimize the lower bound

$$\log p(x) = \log \sum_z q(z | x) \frac{p(x, z)}{q(z | x)} \geq \sum_z q(z | x) \log \frac{p(x, z)}{q(z | x)}, \quad (3)$$

which allows for tractable Monte Carlo approximation.¹

Empirically, we find directly optimizing this lower bound (3) to be more difficult than optimizing the marginal likelihood (1), often requiring a baseline:

$$\begin{aligned} & \nabla_q \sum_z q(z | x) \log \frac{p(x, z)}{q(z | x)} \\ &= \nabla_q \sum_z q(z | x) \log p(x | z) - KL[q(z | x) || p(z)] \\ &= \nabla_q \sum_z q(z | x) (\log p(x | z) - B) - KL[q(z | x) || p(z)], \end{aligned} \quad (4)$$

where the baseline B is not a function of z .² Computing this baseline B can be expensive, potentially requiring multiple evaluations of $p(x | z)$. A common choice of baseline is the sample-average baseline, where $B = \sum_{z' \in Z} \log p(x | z')$ and Z is a set of iid samples.³ Additionally, even in scenarios where exact marginalization is tractable, prior work has included a baseline. This begs the questions: Is a baseline necessary, and why does optimizing the marginal likelihood not require a baseline?

We show that the gradient of the marginal likelihood (2) already contains a built-in baseline, and use that to derive a simple and computable lower bound of the marginal likelihood (1) whose gradient does not require the manual engineering of a baseline.

¹The gap between the two is given by $KL[q(z | x) || p(z | x)]$.

²The derivative is a linear operator, and $\nabla \sum_z q(z | x) B = B \cdot \nabla \sum_z q(z | x) = B \cdot 0$.

³A more efficient alternative is the leave-one-out baseline, which is efficient if the results of $p(x | z)$ are already available for a set of iid z . This can be interpreted as the advantage over a chosen z over the average of the alternatives.

2 Logsumexp

2.1 The gradient of logsumexp approximates the exponentiated regret

We start by reviewing the properties of the gradient of the logsumexp function, before proposing an efficient estimator for variational learning.

The key component of the marginal likelihood is the logsumexp operation, a smooth version of max: $\log \sum_z \exp f(z)$.⁴ Similar to the gradient of the marginal likelihood, the gradient of logsumexp is given by

$$\nabla \log \sum_{z' \in Z} \exp f(z') = \frac{e^{f(z)}}{\sum_{z' \in Z} e^{f(z')}} \nabla f(z) = e^{f(z) - \log \sum_{z' \in Z} e^{f(z')}} \nabla f(z) \approx \underbrace{e^{f(z) - \max_{z' \in Z} f(z')}}_R \nabla f(z).$$

The last approximation holds because logsumexp is a smooth approximation of max. We can therefore think of $\log \sum_{z' \in Z} \exp f(z') \approx \max_{z' \in Z} f(z)$ as a built-in baseline. R is also the exponentiated regret, the difference between a given $f(z)$ and the best $f(z)$. Since the regret is non-positive, the exponentiated regret is always between 0 and 1.

2.2 Approximating the gradient of the marginal likelihood

Given these properties of the gradient of logsumexp, we return to variational learning.

In the marginal likelihood setting (equation 1) $f(z) = \log p(x, z)$, and the regret compares a given $\log p(x, z)$ to the best $\log p(x, z^*)$. In the variational setting (equation (3)), the gradient of the ELBO does not have a built-in baseline or a regret interpretation.

We propose to approximate the gradient of logsumexp by using a restriction of logsumexp to $Z' \subseteq Z$:

$$\log \sum_{z \in Z'} \exp \log p(x, z). \quad (5)$$

We learn a proposal distribution $q(z | x)$ in order to choose Z' .

We instead optimize q to approximate $p(z | x)$ by optimizing

$$\log \sum_z \exp \log p(x, z) - KL[p(z | x) || q(z | x)] \approx \underbrace{\log \sum_{z \in Z'} \exp \log \tilde{p}(x, z)}_{\text{ML}} - \underbrace{\sum_{z \in Z'} \tilde{p}(z | x) \log \frac{\tilde{p}(z | x)}{\tilde{q}(z | x)}}_{\text{KL}}, \quad (6)$$

where $\tilde{p}(z | x) = \frac{p(x, z)}{\sum_{z' \in Z'} p(x, z')}$, where the normalizing constant is approximated.⁵

2.3 Gradient estimator analysis

TBD

3 Related work

Contrastive divergence uses MCMC to estimate the gradient of the log partition function. This is identical to differentiating through a Monte Carlo approximation of the log partition function, commonly known as a sampled softmax / contrastive learning.

⁴ $\log \sum \exp(f(z))$ is perhaps better known as the log-partition function.

⁵This should lead to a derivation of contrastive divergence and contrastive wake sleep. This is probably just an autodiff trick for doing contrastive wake sleep.

Wake-sleep and reweighted wake-sleep optimize the evidence plus the forward KL (M-projection), which is our objective. However, the reconstruction term in the wake-sleep objective is identical to the one used in VAEs, which suffers from poor conditioning. This conditioning is improved by leaving the logsumexp operator intact, rather than lower-bounding using Jensen’s inequality.

DiCE is an infinitely differentiable Monte Carlo gradient estimator. It is a surrogate objective that is written to ensure correct higher order derivatives. Our surrogate objective is also designed to ensure that the gradient holds certain properties, but we focus on controlling the conditioning of the gradient rather than on the correctness of higher order derivatives.